

Project 2

Richard Affolter

4 3 2021

Problem 4: Responsibilities and prior for biased coins

$$\begin{aligned} P(C_i | D_j) &= \frac{P(C_i, D_j)}{\sum_i P(C_i, D_j)} \\ &= \frac{P(D_j | C_i)P(C_i)}{\sum_i P(D_j | C_i)P(C_i)} \end{aligned}$$

For tossing a coin C_i with probability for heads θ_i $n = 10$ times in total we get

$$P(D_j | C_i) = \binom{n}{k_j} \theta_i^{k_j} \cdot (1 - \theta_i)^{n-k_j}$$

where k_j is the number of heads in trial j .

We can simplify the expression since the binomial coefficient does not depend on i and we know that in this case the prior $P(C_i) = \lambda, i = 1, 2$ for both coins is 0.5.

$$\begin{aligned} &\frac{P(D_j | C_i)P(C_i)}{\sum_i P(D_j | C_i)P(C_i)} \\ &= \frac{\binom{n}{k_j} \theta_i^{k_j} \cdot (1 - \theta_i)^{n-k_j} \lambda}{\sum_i \binom{n}{k_j} \theta_i^{k_j} \cdot (1 - \theta_i)^{n-k_j} \lambda} \\ &= \frac{\theta_i^{k_j} \cdot (1 - \theta_i)^{n-k_j}}{\sum_i \theta_i^{k_j} \cdot (1 - \theta_i)^{n-k_j}} \end{aligned}$$

$$\gamma_{1,1} = P(C_1 | D_1) = \frac{0.6^4 \cdot 0.4^6}{0.6^4 \cdot 0.4^6 + 0.4^4 \cdot 0.6^6} = 0.308$$

$$\begin{aligned} \gamma_{1,2} = P(C_2 | D_1) &= \frac{0.4^4 \cdot 0.6^6}{0.4^4 \cdot 0.6^6 + 0.6^4 \cdot 0.4^6} = 0.692 \\ &= 1 - P(C_1 | D_1) \end{aligned}$$

$$\gamma_{2,1} = P(C_1 | D_2) = \frac{0.6^8 \cdot 0.4^2}{0.6^8 \cdot 0.4^2 + 0.4^8 \cdot 0.6^2} = 0.919$$

$$\begin{aligned} \gamma_{2,2} = P(C_2 | D_2) &= \frac{0.4^8 \cdot 0.6^2}{0.4^8 \cdot 0.6^2 + 0.6^8 \cdot 0.4^2} = 0.081 \\ &= 1 - P(C_1 | D_2) \end{aligned}$$

To update the mixture weights:

$$\hat{\lambda}_k = \frac{1}{N} \sum_{i=1}^N \gamma_{i,k}$$
$$P(C_1) = \hat{\lambda}_1 = \frac{1}{2}(0.308 + 0.919) = 0.6135$$
$$P(C_2) = \hat{\lambda}_2 = \frac{1}{2}(0.692 + 0.081) = 0.3865$$
$$= 1 - P(C_1)$$

Data Analysis Problem 5: Learning a mixture model for two biased coins

(a) Load the data stored in the file ‘cointoss.csv’

```
data = read.csv(file = "cointoss.csv", header = TRUE)

# count the number of heads for each trial
trials = colSums(data)

# count the number of tosses per trial
L = dim(data)[1]

# count the number of trials
N = dim(data)[2]
```

(b) Start with random priors (mixture weights) (e.g., $\lambda = (0.5, 0.5)$), and probabilities for heads (= 1) for each of the coins.

```
#initialize the weights and the probabilities for heads
lambda = c(0.5,0.5)
theta = c(0.6,0.5)
```

(c) E-step: use the prior and coin probabilities to compute the joint probability of the observed and hidden data. Compute the responsibilities γ and the log of the observed likelihood.

```
# P(C_i,D_j) meaning coin i and trial j
get_joint_prob = function(i,j){
  return (lambda[i] * dbinom(trials[j],L,theta[i]))
}

get_log_cond_prob = function(i,j){
  return (log(dbinom(trials[j],L,theta[i])))
}

# fill a matrix (2xN) with all joint probabilities
joint_prob = outer(1:2,1:N, FUN = get_joint_prob)
```

```

# fill a matrix (2xN) with all log conditional probabilities
log_cond_prob = outer(1:2, 1:N, FUN = get_log_cond_prob)

# compute the marginal of D_j
marginal = colSums(joint_prob)

# compute the responsibility of component i for observation j
get_responsibility = function(i,j){
  return (joint_prob[i,j]/marginal[j])
}

# fill a matrix (2xN) with all responsibilites
responsibilites = outer(1:2, 1:N, FUN = Vectorize(get_responsibility))

# compute the log of the observed likelihood like on slide 18
part1 = sum(rowSums(responsibilites)*log(lambda)) # first part of the sum
part2 = sum(responsibilites * log_cond_prob) # second part of the sum
l_obs = part1 + part2

```

(d) M-step: use the responsibilities to recompute the priors and the probability of heads (1) of each coin.

```

# maximize w.r.t. lambda
lambda = rowSums(responsibilites)/N

# maximize w.r.t. theta
temp = (responsibilites %*% trials)/
  (responsibilites %*% trials + responsibilites %*% (L-trials))
theta = drop(temp) #get rid of the unneeded dimensions

```

(e) Reiterate over E- and M-step until convergence of the likelihood.

With the functions defined above:

```

EM = function(data, theta, lambda){

  # count the number of heads for each trial
  trials = colSums(data)

  # count the number of tosses per trial
  L = dim(data)[1]

  # count the number of trials
  N = dim(data)[2]

  # initialize log likelihood of previous iteration
  old_l_obs = 0

  iteration = 0

```

```

progress = 1

# loop while we do progress until max number of iterations
while(progress > 1e-7 && iteration < 1e6){

#----- E-step -----

#  $P(C_i, D_j)$  meaning coin  $i$  and trial  $j$ 
get_joint_prob = function(i,j){
  return (lambda[i] * dbinom(trials[j],L,theta[i]))
}

get_log_cond_prob = function(i,j){
  return (log(dbinom(trials[j],L,theta[i])))
}

# fill a matrix (2xN) with all joint probabilities
joint_prob = outer(1:2,1:N, FUN = get_joint_prob)

# fill a matrix (2xN) with all log conditional probabilities
log_cond_prob = outer(1:2,1:N, FUN = get_log_cond_prob)

# compute the marginal of  $D_j$ 
marginal = colSums(joint_prob)

# compute the responsibility of component  $i$  for observation  $j$ 
get_responsibility = function(i,j){
  return (joint_prob[i,j]/marginal[j])
}

# fill a matrix (2xN) with all responsibilities
responsibilites = outer(1:2, 1:N, FUN = Vectorize(get_responsibility))

# compute the log of the observed likelihood like on slide 18
part1 = sum(rowSums(responsibilites)*log(lambda)) # first part of the sum
part2 = sum(responsibilites * log_cond_prob) # second part of the sum
l_obs = part1 + part2

#----- M-step -----

# maximize w.r.t. lambda
lambda = rowSums(responsibilites)/N

# maximize w.r.t. theta
temp = (responsibilites %*% trials)/
  (responsibilites %*% trials + responsibilites %*% (L-trials))
theta = drop(temp) #get rid of the unneeded dimensions

#----- update conditions for loop termination-----
progress = abs(l_obs - old_l_obs)
old_l_obs = l_obs
iteration = iteration + 1
}

```

```

    return(list(lambda = lambda, theta = theta, gamma = responsibilites,
               log_likelihood = l_obs, progress=progress, iteration = iteration))
}

```

We can start the EM several times independently with different priors and choose the result with the highest likelihood

```

set.seed(123)

best_likelihood = -Inf

# run the algorithm 100 times
for(i in 1:100){
  # generate random initialization for theta and lambda
  random = runif(2)
  theta_init = c(random[1], 1-random[1])
  lambda_init = c(random[2], 1-random[2])

  current_result = EM(data = data, theta = theta_init, lambda = lambda_init)

  if(current_result$log_likelihood > best_likelihood){
    best_result = current_result
    best_likelihood = current_result$log_likelihood
  }
}

```

(f) Print the probability of heads for each coin and the mixture weights λ . Plot a heatmap of the responsibilities. How many trials belong to each coin?

```

for(i in 1:2){
  cat(
    paste0("Probability of heads for coin", i, ": ",
          round(best_result$theta[i], digits = 6), "\n") )
}

```

```

## Probability of heads for coin1: 0.451041
## Probability of heads for coin2: 0.572846

```

```

cat("\n Mixture weights: (", round(best_result$lambda, digits = 6), ",)")

```

```

##
## Mixture weights: ( 0.841068 0.158932 )

```


So we have $\theta_1 = 0.451041$ and $\theta_2 = 0.572846$.

The mixture weights are $\lambda_1 = 0.841068$ and $\lambda_2 = 1 - \lambda_1 = 0.158932$

```

library(ComplexHeatmap)
rownames(best_result$gamma) = paste("coin",1:2)
colnames(best_result$gamma) = rep('',N)
for(i in seq(0,N,25)){
  colnames(best_result$gamma)[i] = i
}
colnames(best_result$gamma)[1]=1
Heatmap(best_result$gamma,
  #col = RColorBrewer::brewer.pal(name = "OrRd", n = 9),
  #col = colorRampPalette(c("black", "white"),10),
  #width = unit(10, "cm"),
  cluster_rows = FALSE,
  cluster_columns = FALSE,
  column_title = "trials",
  column_title_side = "bottom",
  #row_title = "coins",
  name = "responsibilities")

```



```

cat(sum(best_result$gamma[1,]>0.5),
  "trials belong to coin 1 and",
  sum(best_result$gamma[2,]>0.5),
  "trials belong to coin 2.")

```

```
## 171 trials belong to coin 1 and 29 trials belong to coin 2.
```